

# High Performance Scientific computing

## Lecture 3

S. Gopalakrishnan

# Von Neumann Architecture




Image: Wikipedia

# Program Code Levels

- High-level language
  - Level of abstraction closer to problem domain
  - Provides for productivity and portability

```
for( i=0; i<=NR+1; i++ )      /* Copy the values into told */  
    for( j=0; j<=NC+1; j++ )  
        told[i][j] = t[i][j];
```

# Program Code Levels

- High-level language

- Level of abstraction closer to problem domain
- Provides for productivity and portability

```
for( i=0; i<=NR+1; i++ )      /* Copy the values into told */  
    for( j=0; j<=NC+1; j++ )  
        told[i][j] = t[i][j];
```

- Assembly language

- Textual representation of instructions.
- Compilation and linking of program will yield this

```
je else  
jmp endif  
else:  
endif:  
    sarl $1, %edx  
    movl %edx, %eax  
    addl %eax, %edx  
    addl %eax, %edx  
    addl $1, %edx  
    addl $1, %ecx
```

- Hardware representation

# Program Code Levels

- High-level language
  - Level of abstraction closer to problem domain
  - Provides for productivity and portability

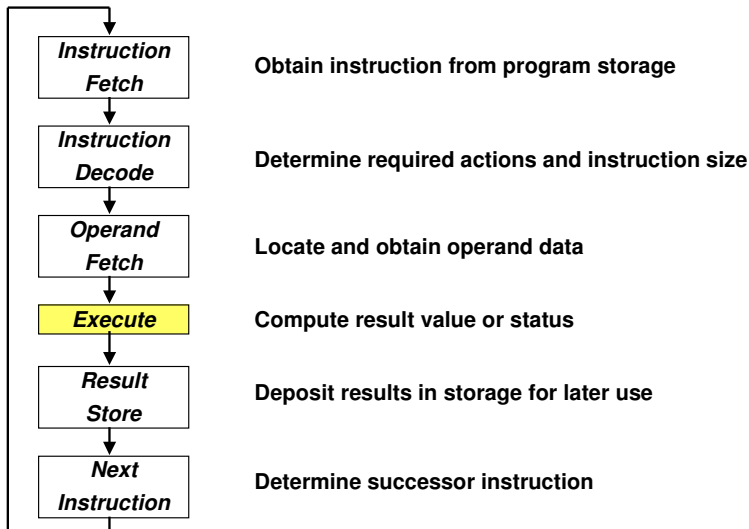
```
for( i=0; i<=NR+1; i++ )      /* Copy the values into told */  
    for( j=0; j<=NC+1; j++ )  
        told[i][j] = t[i][j];
```

- Assembly language
  - Textual representation of instructions.
  - Compilation and linking of program will yield this

```
je else  
jmp endif  
else:  
endif:  
    sarl $1, %edx  
    movl %edx, %eax  
    addl %eax, %edx  
    addl %eax, %edx  
    addl $1, %edx  
    addl $1, %ecx
```

- Hardware representation
  - Binary digits (bits)
  - Encoded instructions and data

## Execution cycle



Credit: Based on notes from EC232 UMass – Amherst

## Instruction Pipelining

Increases instruction throughput, i.e number of instructions executed per second. But time for each instruction remains the same.

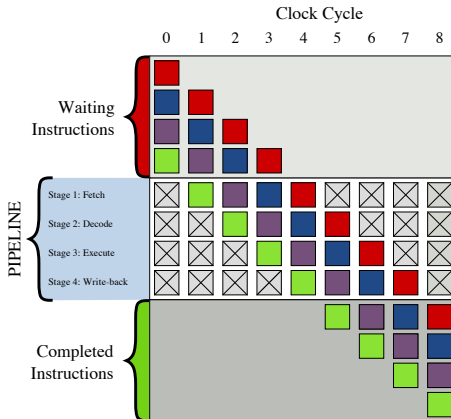


Image: Wikipedia

## Pipelining Issues

Data dependencies can result in slowing of throughput.

```
for( i=0; i<=LAST; i++ )  
  a[i]=b[i]+c[i]
```

No deadlock in previous sample, pipelining will work like a charm.



## Pipelining Issues

Data dependencies can result in slowing of throughput.

```
for( i=0; i<=LAST; i++ )  
a[i]=b[i]+c[i]
```

No deadlock in previous sample, pipelining will work like a charm.

```
for( i=0; i<=LAST; i++ )  
a[i]=b[i]+a[i-1]
```

This is a deadlock situation since  $a[i]$  depends on  $a[i - 1]$

## Pipelining Issues

Bubbles in pipelines due to dependencies.

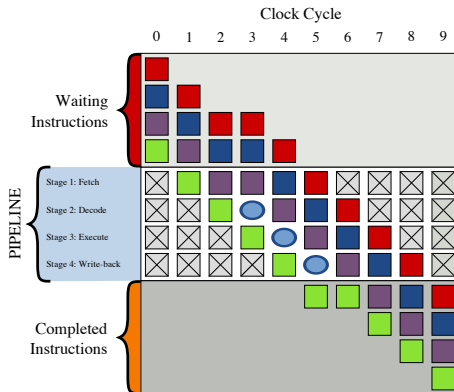
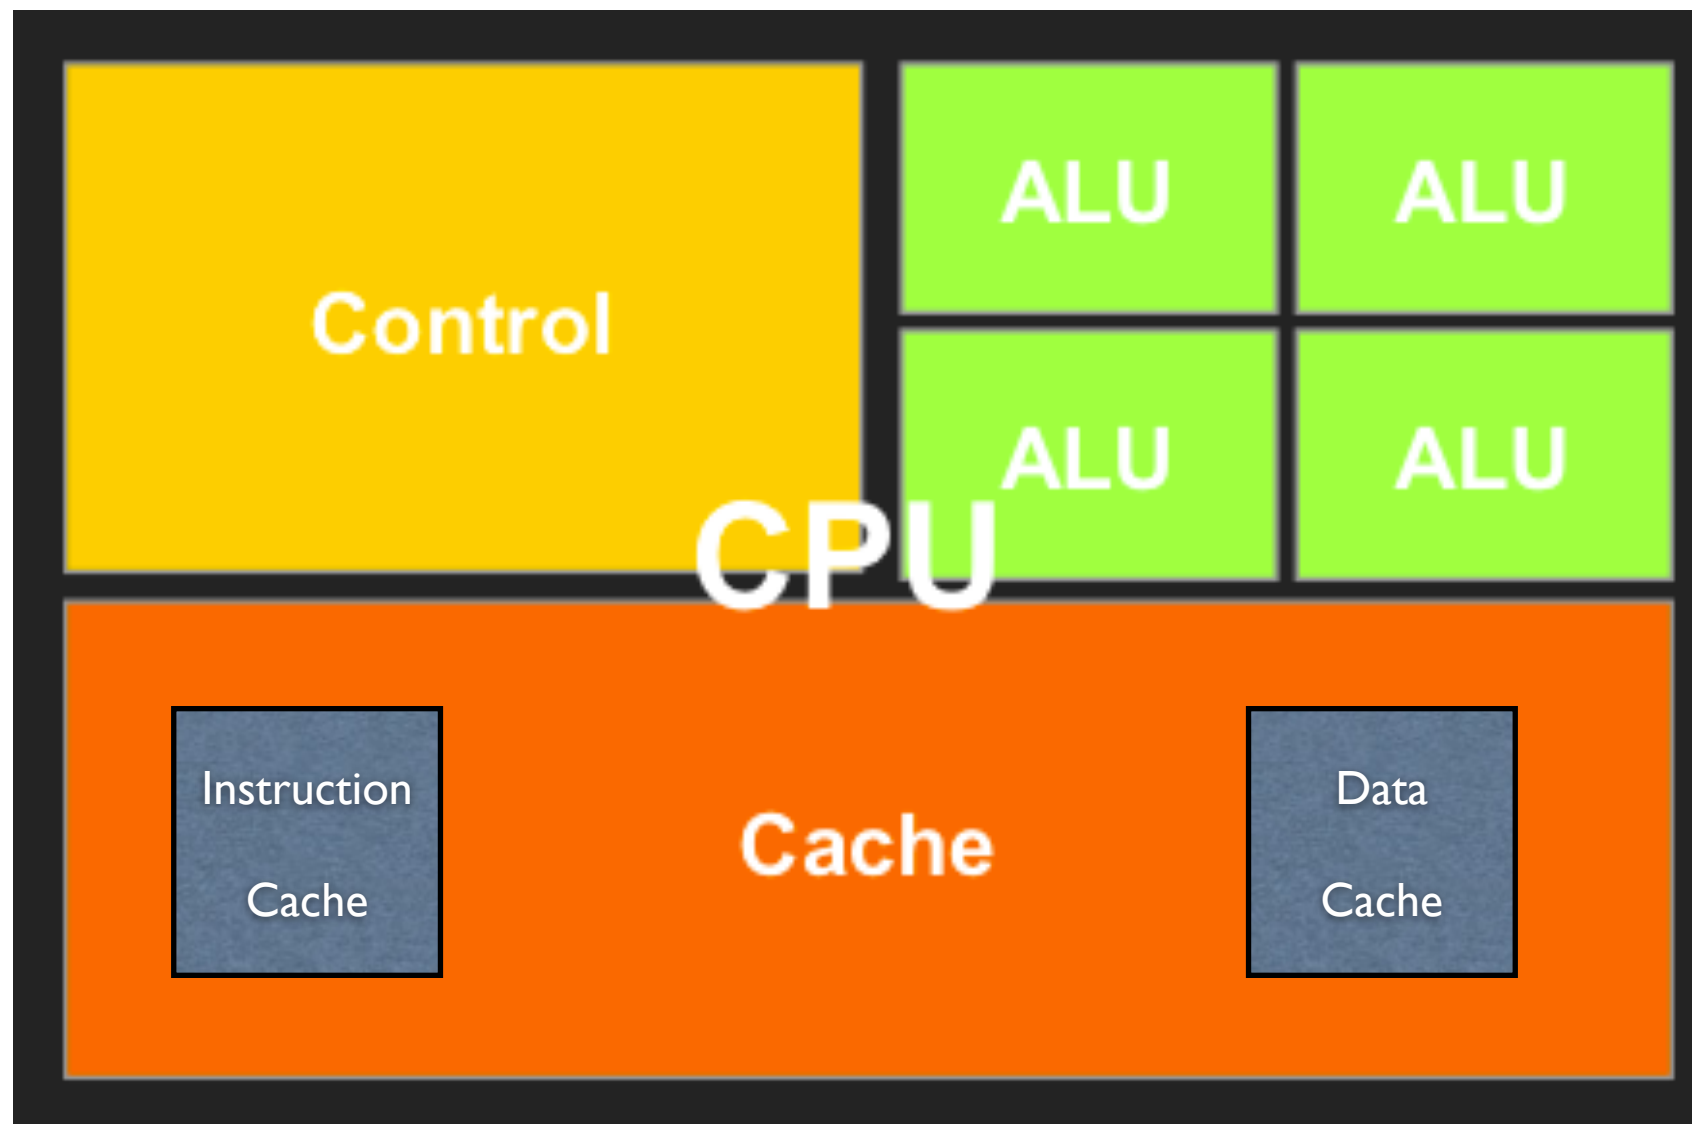


Image: Wikipedia

# Memory Issues

# CPU Block Diagram



# Memory Basics

Users want large and fast memories!

Fact

- Large memories are slow
- Fast memories are small


Assume CPU runs at 3GHz

Every instruction requires 4B of instruction and at least one memory access (4B of data)

- $3 * 8 = 24\text{GB/sec}$

Memory bandwidth and access time is a performance bottleneck

## Processor-DRAM Memory Performance Gap Motivation for Memory Hierarchy



# Memory Basics

Small memories are fast

--So just write small programs

“640 K of memory should be enough for anybody” -- Bill Gates, 1981

Today's programs require large memories

- Powerpoint 2003 – 25 megabytes
- Scientific applications may require Gigabytes of memory

How do we create a memory that is large, cheap and fast (most of the time)?

Strategy: Provide a Small, Fast Memory which holds a subset of the main memory – called **cache**

- Keep frequently-accessed locations in fast cache
- Cache retrieves more than one word at a time
- Sequential accesses are faster after first access

# Memory Basics

## Hierarchy of Levels

- Uses smaller and faster memory technologies close to the processor
- Fast access time in highest level of hierarchy
- Cheap, slow memory furthest from processor
- The aim of memory hierarchy design is to have access time close to the highest level and size equal to the lowest level

## **Basic Philosophy**

Move data into ‘smaller, faster’ memory

Operate on it

Move it back to ‘larger, cheaper’ memory


- How do we keep track if changed


What if we run out of space in ‘smaller, faster’ memory?

Important Concepts: Latency, Bandwidth



# Memory hierarchy






- Notice that the data width is changing
  - Why?
- Bandwidth: Transfer rate between various levels
  - CPU-Cache: 24 GBps
  - Cache-Main: 0.5-6.4GBps
  - Main-Disk: 187MBps (serial ATA/1500)

# Virtual Memory and Paging

Virtual memory  
(per process)

Physical  
memory




RAM


Disk

# Memory Hierarchy Terminology

- **Hit**: data appears in upper level in block X
- **Hit Rate**: the fraction of memory accesses found in the upper level
- **Miss**: data needs to be retrieved from a block in the lower level (Block Y)
- **Miss Rate** =  $1 - (\text{Hit Rate})$
- **Hit Time**: Time to access the upper level which consists of Time to determine hit/miss + upper level access time
- **Miss Penalty**: Time to replace a block in the upper level + Time to deliver the block to the processor
- Note: **Hit Time**  $\ll$  **Miss Penalty**



# Current Memory Hierarchy



Speed(ns):	1ns	2ns	6ns	100ns	10,000,000ns
Size (MB):	0.0005	0.1	1-4	1000-6000	500,000
Cost (\$/MB):	--	\$10	\$3	\$0.01	\$0.002
Technology:	Regs	SRAM	SRAM	DRAM	Disk

- **Cache - Main memory: Speed**
- **Main memory – Disk (virtual memory): Capacity**